

Enhanced Logistic Regression on the Titanic Dataset

Soorya s

October 3, 2025

1 Introduction

This report presents a logistic regression model applied to the Titanic dataset. The model includes:

- L_2 regularization to control overfitting,
- momentum gradient descent to smooth convergence,
- learning rate decay for stability,
- grid search for hyperparameter tuning.

We preprocess the data, derive the mathematics, implement in Python, and present the results.

2 Mathematical Formulation

Given $X \in \mathbb{R}^{m \times n}$ (feature matrix), $\vec{w} \in \mathbb{R}^n$ (weights), and $b \in \mathbb{R}$ (bias):

$$\vec{z} = X\vec{w} + b\mathbf{1}, \quad \hat{p} = \sigma(\vec{z}) = \frac{1}{1 + e^{-\vec{z}}}.$$

The loss function is binary cross-entropy with L_2 penalty:

$$L(\vec{w}, b) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \ln \hat{p}^{(i)} + (1 - y^{(i)}) \ln(1 - \hat{p}^{(i)}) \right] + \frac{\lambda}{2m} \|\vec{w}\|^2.$$

Gradients are:

$$\nabla_{\vec{w}} L = \frac{1}{m} X^\top (\hat{\vec{p}} - \vec{y}) + \frac{\lambda}{m} \vec{w}, \quad \frac{\partial L}{\partial b} = \frac{1}{m} \sum_{i=1}^m (\hat{p}^{(i)} - y^{(i)}).$$

Momentum Update

Momentum accumulates gradients:

$$\vec{v}_t = \beta \vec{v}_{t-1} + (1 - \beta) \nabla_{\vec{w}} L, \quad \vec{w} \leftarrow \vec{w} - \alpha_t \vec{v}_t, \quad b \leftarrow b - \alpha_t v_{b,t}.$$

Learning Rate Decay

We use reciprocal decay:

$$\alpha_t = \frac{\alpha_0}{1 + kt}.$$

3 Preprocessing

We use z-score normalization:

$$x_j^{(i)} = \frac{x_j^{(i)} - \mu_j}{\sigma_j}, \quad \sigma_j = 0 \Rightarrow \sigma_j = 1.$$

The dataset is split:

Train: 80%, Validation: 10%, Test: 10%.

4 Implementation

Data Loading and Preprocessing

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.metrics import confusion_matrix
6
7 # Load Titanic dataset
8 df = pd.read_csv("train.csv")
9 df = df[["Survived", "Pclass", "Sex", "Age", "SibSp", "Parch", "Fare"]]
10
11 # Handle missing values
12 df["Age"].fillna(df["Age"].mean(), inplace=True)
13
14 # Encode categorical variable
15 df["Sex"] = (df["Sex"] == "male").astype(int)
16
17 # Build feature matrix and target
18 X_full = df[["Pclass", "Sex", "Age", "SibSp", "Parch", "Fare"]].
19     to_numpy(float)
20 y_full = df["Survived"].to_numpy(float)
21 m_full, n = X_full.shape
22
23 # Split 80/10/10
24 rng = np.random.default_rng(42)
25 perm = rng.permutation(m_full)
26 m_train = int(0.80 * m_full)
27 m_val = int(0.10 * m_full)
28 idx_train = perm[:m_train]
29 idx_val = perm[m_train:m_train+m_val]
30 idx_test = perm[m_train+m_val:]
31 X_train, y_train = X_full[idx_train], y_full[idx_train]
32 X_val, y_val = X_full[idx_val], y_full[idx_val]
33 X_test, y_test = X_full[idx_test], y_full[idx_test]
```

```

34 # Normalize
35 mu = X_train.mean(axis=0)
36 sigma = X_train.std(axis=0, ddof=0)
37 sigma[sigma == 0] = 1.0
38 X_train = (X_train - mu) / sigma
39 X_val = (X_val - mu) / sigma
40 X_test = (X_test - mu) / sigma

```

Model and Loss

```

1 def sigmoid(z):
2     z = np.clip(z, -500, 500)
3     return 1 / (1 + np.exp(-z))
4
5 def f(X, w, b):
6     return sigmoid(X @ w + b)
7
8 def loss(X, y, w, b, lam):
9     m = len(y)
10    p = f(X, w, b)
11    eps = 1e-12
12    p = np.clip(p, eps, 1 - eps)
13    ce = -np.mean(y*np.log(p) + (1-y)*np.log(1-p))
14    l2 = (lam / (2*m)) * np.sum(w*w)
15    return ce + l2

```

Momentum Step and LR Decay

```

1 def step_mom(X, y, w, b, v_w, v_b, lr, lam, beta=0.9):
2     m = len(y)
3     p = f(X, w, b)
4     e = p - y
5     g_w = (X.T @ e) / m + (lam/m) * w
6     g_b = e.mean()
7
8     v_w = beta * v_w + (1-beta) * g_w
9     v_b = beta * v_b + (1-beta) * g_b
10
11    w = w - lr * v_w
12    b = b - lr * v_b
13    return w, b, v_w, v_b
14
15 def lr_decay(alpha0, t, k=1e-3):
16     return alpha0 / (1 + k*t)
17
18 def accuracy(y, p, thr=0.5):
19     yhat = (p >= thr).astype(float)
20     return (yhat == y).mean()

```

Grid Search

```

1 lam_vals = [0.0, 0.005, 0.01, 0.05, 0.1]

```

```

2 alpha0_vals = [0.05, 0.1, 0.2, 0.3]
3 beta_vals = [0.85, 0.9, 0.95]
4
5 best_val = 0
6 best = {}
7
8 for lam in lam_vals:
9     for a0 in alpha0_vals:
10        for beta in beta_vals:
11            w = np.zeros(n); b = 0.0
12            v_w = np.zeros(n); v_b = 0.0
13            prev_L = None
14            for epoch in range(1, 3001):
15                lr = lr_decay(a0, epoch, k=1e-3)
16                w, b, v_w, v_b = step_mom(X_train, y_train, w, b, v_w,
17                                         v_b, lr, lam, beta)
18                L = loss(X_train, y_train, w, b, lam)
19                if prev_L is not None and abs(prev_L - L) < 1e-7: break
20                prev_L = L
21            val_acc = accuracy(y_val, f(X_val, w, b))
22            if val_acc > best_val:
23                best_val = val_acc
24                best = {"lam": lam, "a0": a0, "beta": beta}

```

Final Training and Evaluation

```

1 lam, a0, beta = best["lam"], best["a0"], best["beta"]
2 w = np.zeros(n); b = 0.0
3 v_w = np.zeros(n); v_b = 0.0
4
5 for epoch in range(1, 5001):
6     lr = lr_decay(a0, epoch, k=1e-3)
7     w, b, v_w, v_b = step_mom(X_train, y_train, w, b, v_w, v_b, lr, lam
8                               , beta)
9     if epoch % 500 == 0:
10        print("Epoch", epoch, "Train Loss", loss(X_train, y_train, w, b,
11                                                    lam))
12
13 p_tr = f(X_train, w, b)
14 p_va = f(X_val, w, b)
15 p_te = f(X_test, w, b)
16
17 print("Train Acc:", accuracy(y_train, p_tr))
18 print("Val Acc:", accuracy(y_val, p_va))
19 print("Test Acc:", accuracy(y_test, p_te))

```

5 Results

Grid search gave best parameters: $\lambda = 0$, $\alpha_0 = 0.1$, $\beta = 0.85$. Final accuracies:

Train: $\approx 82.9\%$, Validation: $\approx 77.5\%$, Test: $\approx 78.9\%$.

6 Conclusion

The logistic regression model with momentum and learning rate decay converges smoothly and avoids oscillations. Hyperparameter tuning confirmed that regularization was not needed, since the dataset was large enough and momentum already stabilized training.